

A support patch of $z_j^{(t)}$ contains visual evidence of $z_j^{(t)}$ that **activates** feature j at token t .



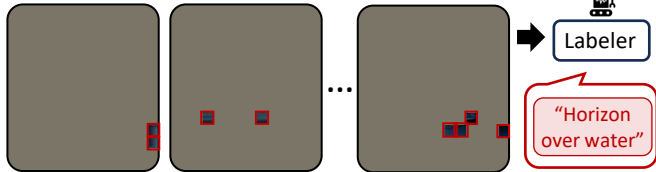
Activation alone is insufficient to explain SAE feature: relevant evidence may lie in their **support patches**.

AL-based (traditional) vs ERF-based (ours) autolabel pipelines

same feature f , same top examples

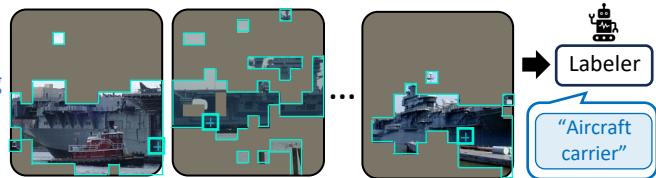
AL-based autolabeling

Show tokens where f fires



ERF-based autolabeling

Show tokens what makes f to fire (ERF)



Masked region



Support patch (sufficient evidence)



Target token (readout site $z_j^{(t)}$)



Activation-only visible patch